

Network Programming

Assignment #2

Design Decisions

Date Submitted:
April 24, 2019

Vedant Patwary (2016A7PS0031P)
Kunal Jain (2016A7PS0022P)

Part 1

Deliverables included: multithread.c, eventdriven.c, client.c, makefile

How to run:

- \$ make
- a. Event-driven server
 - \$./eventdriven
- b. Threaded server
 - \$./threadedserver
- c. Client
 - \$./client <N> <M> <T>

Constraints:

1. Maximum size of message can be 50.
2. Client can have a name of maximum length 40.

Working of the code:

- Run any of the server (./eventdriven or ./threadedserver)
- Then it will prompt the user to enter a port number.
- Enter any free port number and continue.
- To use the server,
 - **Telnet:**
 - Use the command telnet <ip> <port>
 - JOIN <name>
 - LIST
 - UMSG <dest. Name> <msg> : press enter to send
 - BMSG <msg> : press enter to send
 - LEAV
 - **Client:**
 - Run ./client <N> <M> <T> : ex. \$./client 10 10 20
 - The client will print all the tasks it is doing
 - The client will also print the time taken for it to complete the execution
- To exit the server, send SIGINT (Ctrl + C). The server closed successfully.
- To run the server again, the user can use the same port again.

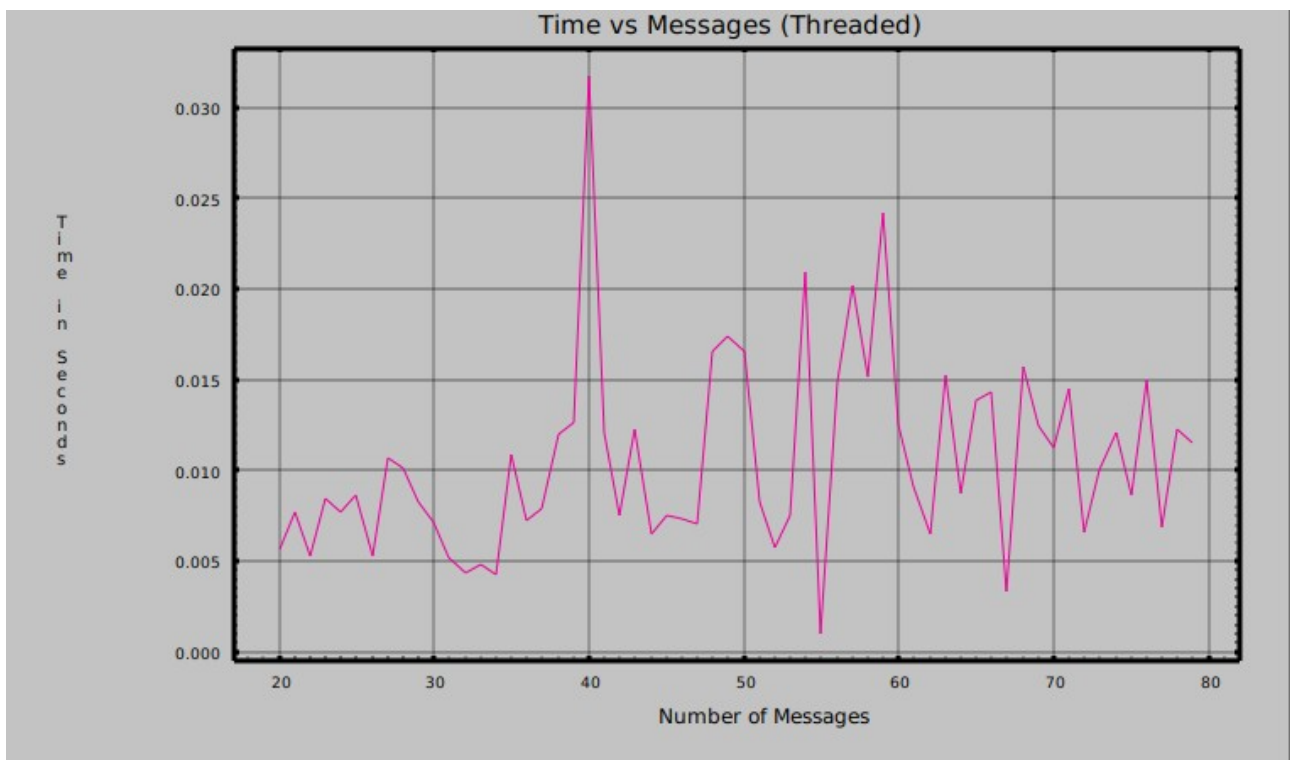
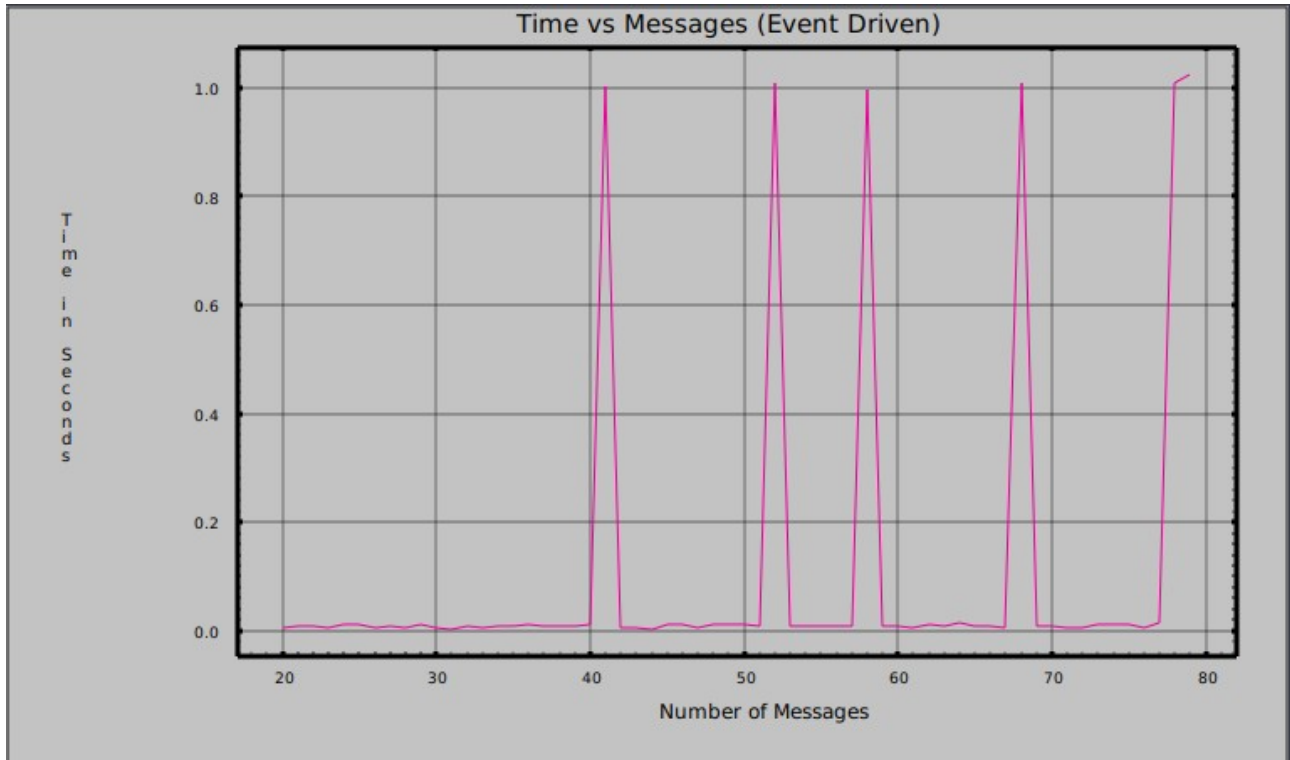
Implementation:

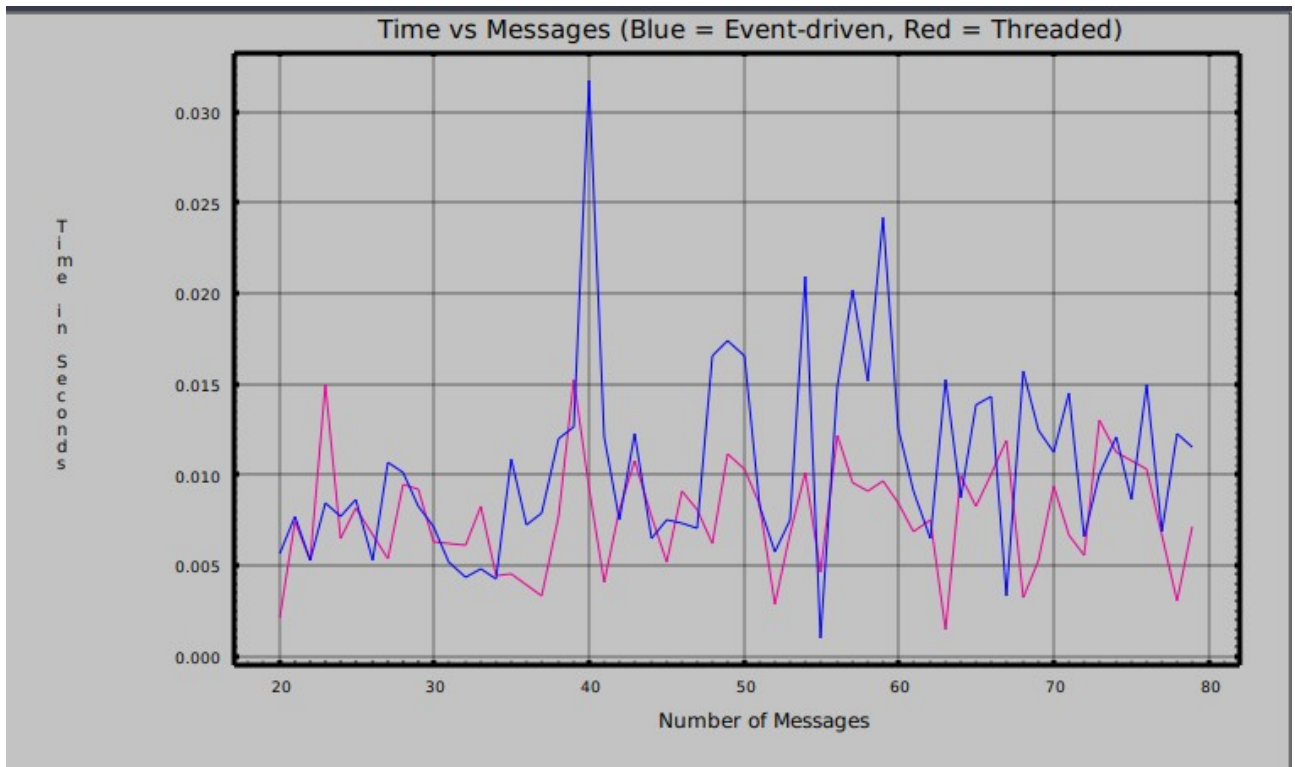
- **Event-driven:**
 - A listening socket is created to accept connections. *epoll_wait* waits for a client to connect to this socket.
 - When listening fd is readable, accept a client into *connfd*.
 - Make the *connfd* non blocking.
 - Add this *connfd* to *epoll*'s read interest list.
 - Whenever client sends some data, this *connfd* becomes readable along with all the other *connfds* (of other clients, if they send some data at the same time).
 - Iterate over the ready list to find out those *connfds* that are ready to be sent to the reading event.
 - Once all the fds are queued into the *reading* event, we will process all the queued events at once. Now these events are queued in the processing event queue.
 - All the events queued in the above mentioned processing queue are now processed and if they're
 - JOIN: Inserting the details of the new client (*connfd* and *client name*) into a globally declared linked list.
 - LIST: Iterates over this linked list and stores list of names of all the clients online.
 - UMSG <dest name> <msg> : Finds the fd of the client with name *dest name* and stores the *msg* and *connfd*.
 - BMSG <msg> : Iterates over all the clients in the global linked list and stores their fds alongwith the *msg*.
 - LEAV: Deletes the entry for the source client from the global linked list.
 - and after storing, it queues the packet to be sent in the writing queue.
 - All the events queued in writing event queue are sent to the desired clients.
- **Threaded**
 - Same data structure is used from event-driven server.
 - We are using the child-thread-accept model.
 - A listening socket is created and *select* waits on this fd to be readable.
 - Whenever this fd is readable, a new client initiated a connection, so a new thread is made.
 - This thread accepts the connection into a *connfd* and uses this fd to communicate with the client.
 - A while(1) loop is used to read the data, process it and write back to client.
 - A conceptual problem arises when accessing the shared global linked list. The server should allow multiple clients to send UMSG, BMSG and LIST but only one client at a time is allowed to use JOIN and LEAV (as they are changing the linked list). This problem is an example of **reader-writer** problem.

- To solve this, we used 3 shared variables:
 - *writing*: is one if someone is writing (using JOIN/LEAV)
 - *writers*: number of threads waiting to change and currently changing the linked list
 - *reading*: number of threads currently reading (using BMSG, UMSG, LIST)
- 1 mutex:
 - *mutex*: to introduce exclusion and synchronization between the shared variables.
- 1 conditional variable:
 - *cond*: to wake the waiting threads
- these all variables and mutexes are used to implement the solution to the problem-
 - The writer waits until there is no thread *reading* or *writing* currently
 - The reader waits if there are any *writers* present
- **Client**
 - Takes N, M, T as CLA
 - Takes port number and then IP address as input.
 - Uses pre-forking to create N processes.
 - A variable maintains a count of children who have spawned to communicate with server. This variable is changed via SIGCHILD signal in parent process.
 - Whenever there is space to spawn, the process connects to the server and starts to communicate
 - send JOIN message alongwith the client number.
 - sends M number of UMSGs to itself
 - send LEAV and then terminates
 - prints the time taken to do all the tasks.

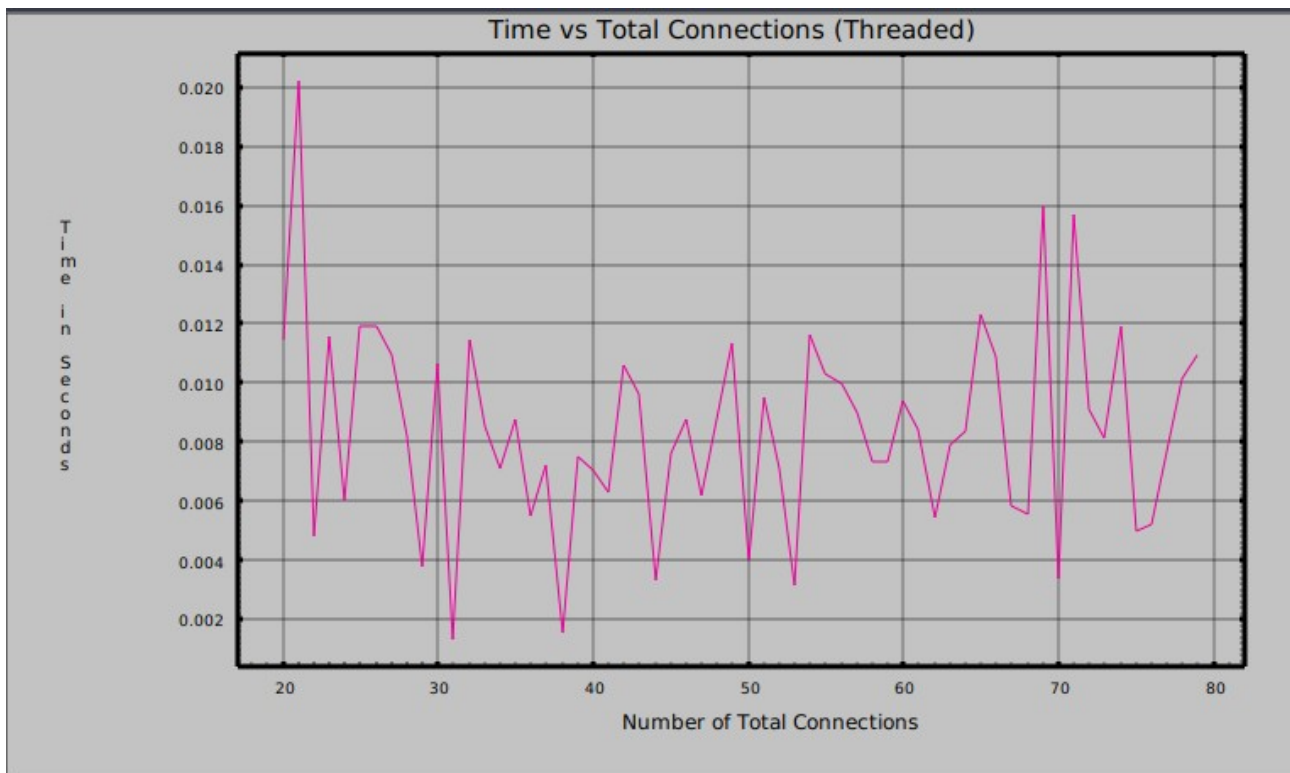
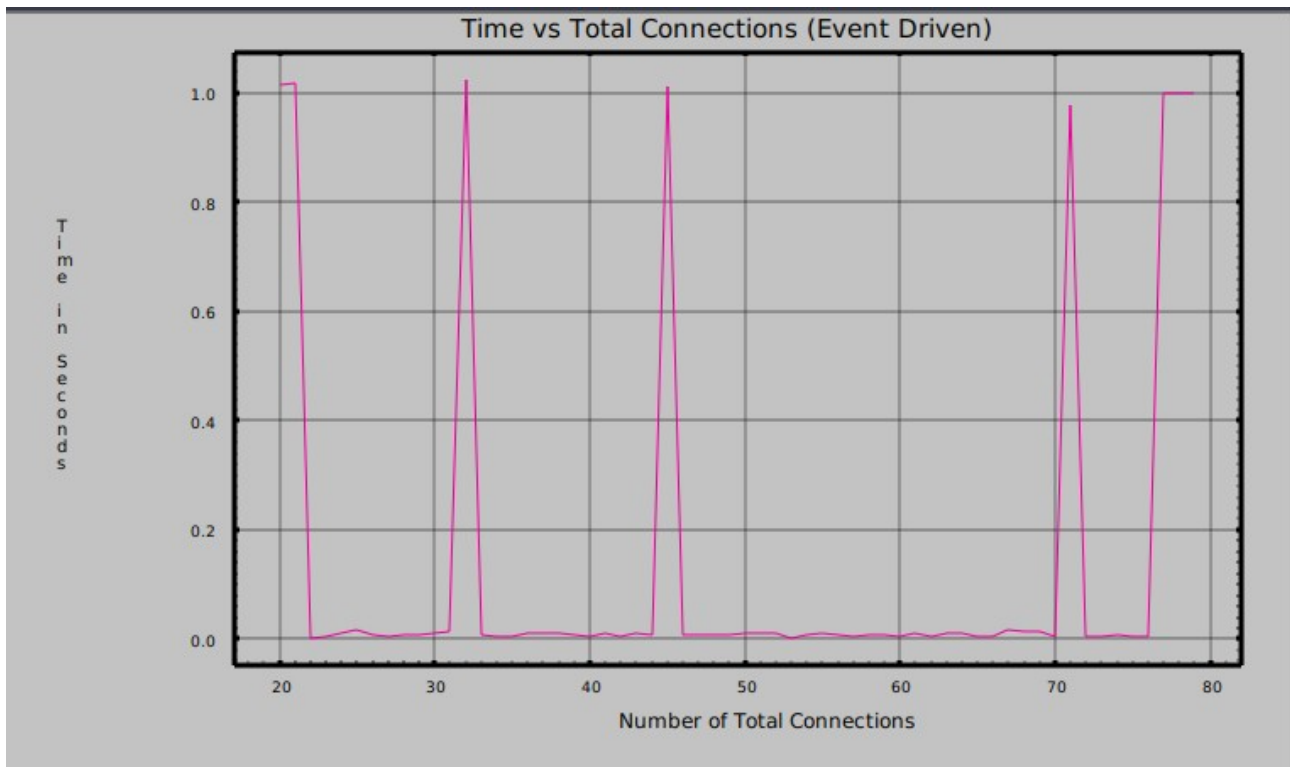
Comparative Analysis:

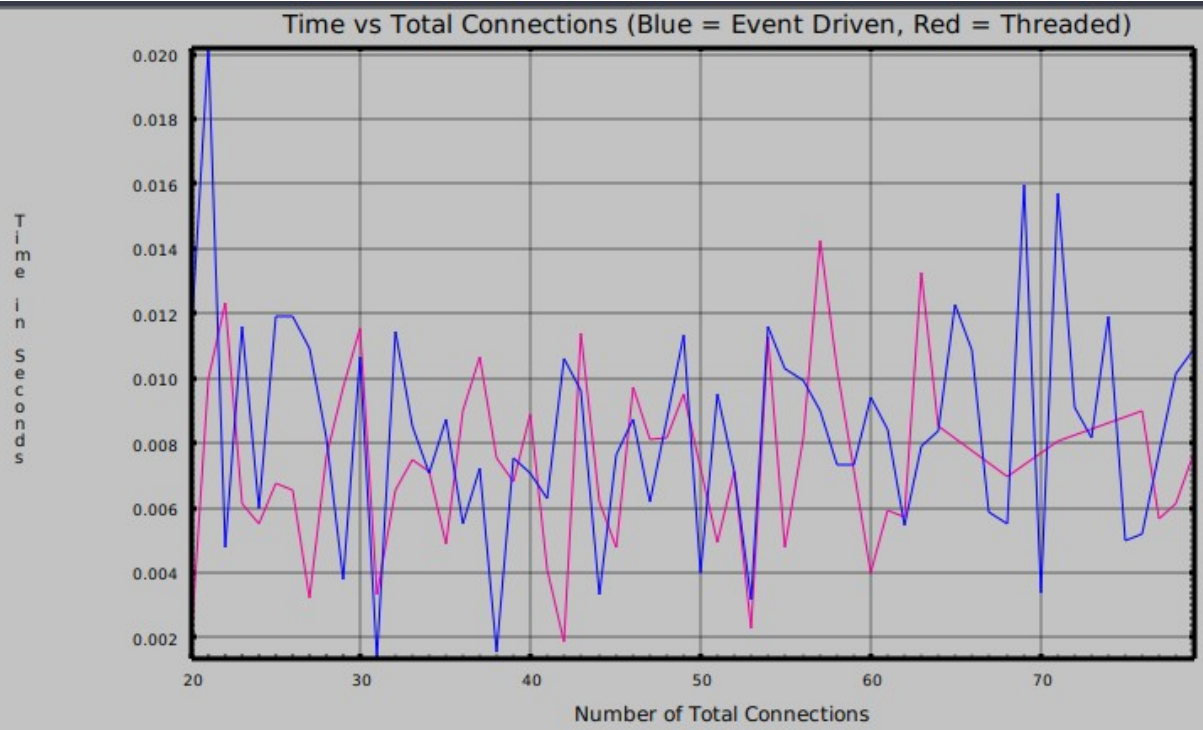
Time vs Number of Messages



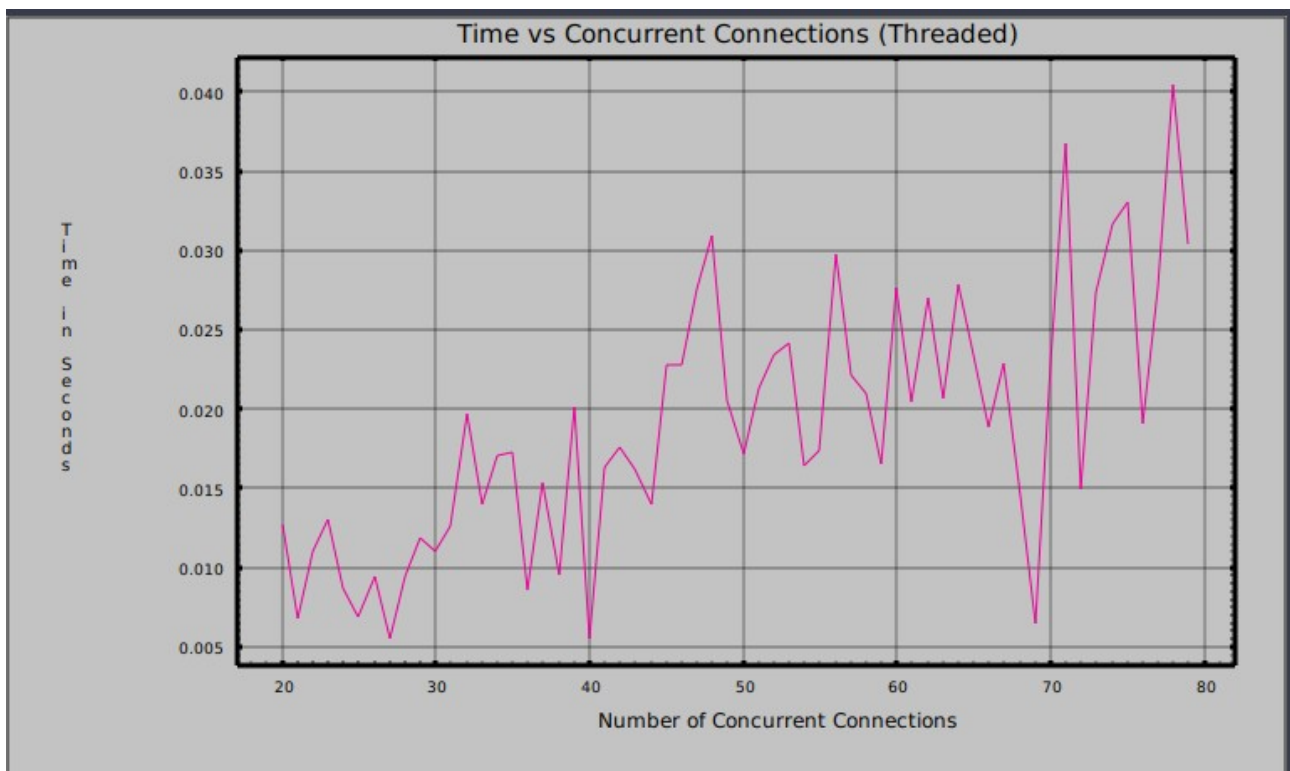
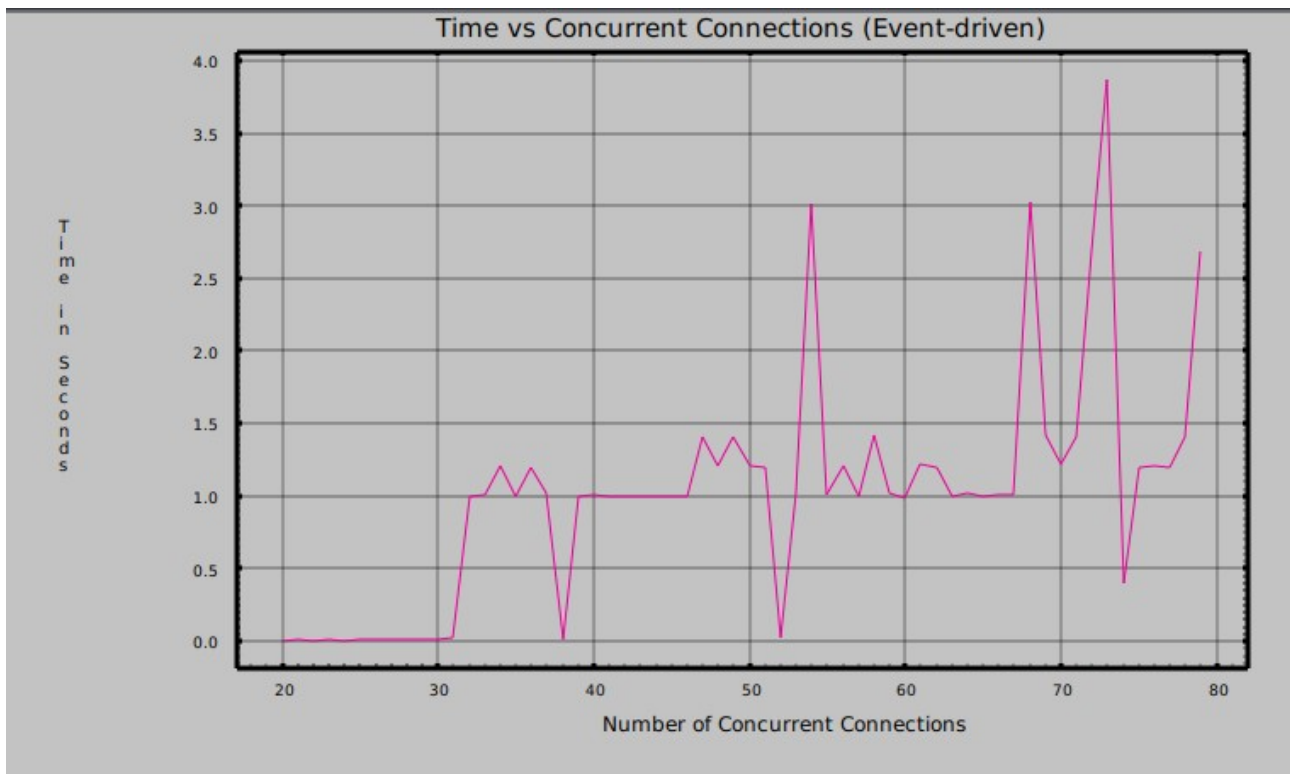


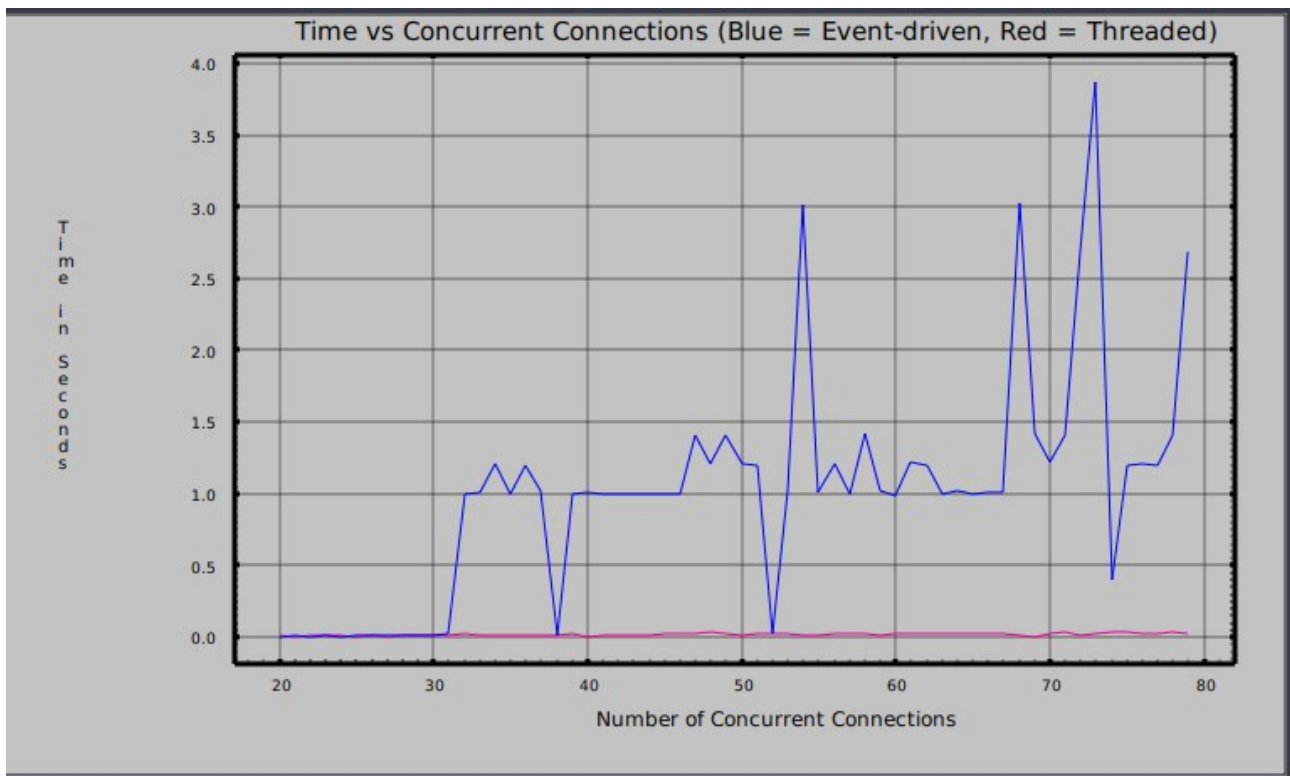
Time vs Total Connections





Time vs Concurrent Connections





Conclusion:

It is evident from the comparison given above that changing the number of concurrent connections has most impact on the time taken to serve the clients. Also, threaded server seems to out-perform event-driven server in case 1 and 3 (not in total connections where both give similar result).

---X---